

IBG Research Computing Cheatsheet

Monitoring Slurm and Job Activity

Monitoring running jobs

View your running jobs

```
squeue -u $USER
```

- Additional fields displayed using the `-l` flag
- Custom output described in man page. E.g., I find `squeue -o "%.12i %.18j %.16q %.8T %.10M %.12l %.24R"` easier to read

View running jobs on a particular qos

```
squeue -q <QOS>
```

Measure performance of completed jobs

The following flags can be combined to suit your needs

Display jobs completed since a particular date

```
sacct -S <MMDD>
```

Display jobs' timing and memory usage

```
sacct -o 'jobid%20,jobname%16,state,elapsed,maxrss'
```

View available nodes and their properties

```
sinfo --Node -o "%.12N %.16P %.11T %.4c %.13C %.8e /%.8m %.30f"
```

- NODELIST is the name of the node, can be specified when submitting via `-nodelist`
- PARTITION is the name of the partition. You likely will never need to use this unless you want a particular partition of the preemptable queue
- STATE - mixed means that some of the CPUS are in use (see below)
- CPUS(A/I/O/T) - on a given node: allocated/idle/other/total nodes
- FREE_MEM / MEMORY free memory / total memory in MB (divide by 1024 to convert to GB)
- AVAIL_FEATURES - architectures/instruction sets requestable via the `--constraint` flag

Example jobs

Basics

Interactive jobs

```
#placeholder
```

Preemptable job

```
#placeholder
```

Job arrays

Job array with a single numeric index

```
#!/bin/bash
#SBATCH --qos=<QOS>
#SBATCH --mem=<MEM>gb
#SBATCH --time=<D-HH:MM> or <HH:MM:SS>
#SBATCH --ntasks=<Number of cpus>
#SBATCH --nodes=1
#SBATCH --array=1-22
#SBATCH -J <jobname>
#SBATCH -o <output dir>/<jobname>_%a
```

```
CHR=${SLURM_ARRAY_TASK_ID}
```

```
ml load <modules>
```

```
program <args> \  
  --input <inputprefix>_"$CHR" \  
  --output <outputprefix>_"$CHR"
```

Job array with a single non-numeric index

This script submit jobs in parallel for jobs with different inputs

```
#!/bin/bash  
...  
#SBATCH --array=0-2  
#SBATCH -J <jobname>  
#SBATCH -o <output dir>/<jobname>_%a
```

```
ii=${SLURM_ARRAY_TASK_ID}
```

```
inputArray=(phenoA phenoB phenoC)
```

```
input=${inputArray[$ii]}
```

```
ml load <modules>
```

```
program <args> \  
  --input $input
```

Job array with a multiple simultaneous indices

This script submit jobs in parallel for jobs with different inputs/outputs

```
#!/bin/bash  
...  
#SBATCH --array=0-2  
#SBATCH -J <jobname>  
#SBATCH -o <output dir>/<jobname>_%a
```

```
ii=${SLURM_ARRAY_TASK_ID}
```

```
inputArray=(phenoA phenoB phenoC)
```

```
outputArray=(outputA outputB outputC)
```

```
input=${inputArray[$ii]}  
output=${outputArray[$ii]}
```

```
ml load <modules>
```

```
program <args> \  
  --input $input \  
  --output $output
```

Job array with a multiple nested non-numeric indices

This script submit jobs in parallel for jobs with arbitrary nested lists arguments (e.g., each model for each phenotype) using integer arithmetic. If this is unfamiliar, you can google “floor division bash” and “modulo bash”.

```
#!/bin/bash  
...  
#SBATCH --array=0-5  
#SBATCH -J <jobname>  
#SBATCH -o <output dir>/<jobname>_%a
```

```
ii=${SLURM_ARRAY_TASK_ID}
```

```
inputArray=(phenoA phenoB phenoC)
modelArray=(modelA modelB)

inputIndex=$(expr $ii % 3)
modelIndex=$(expr $ii / 3)
```

```
input=${inputArray[$inputIndex]}
model=${modelArray[$modelIndex]}
```

```
ml load <modules>
```

```
program <args> \
  --phenotype "$input" \
  --model $model"
```

You can always double check that you didn't screw something up by running a simple loop (in the shell):

```
inputArray=(modelA modelB)
modelArray=(phenoA phenoB phenoC)

for ii in {0..5}
do
  inputArray=(phenoA phenoB phenoC)
  modelArray=(modelA modelB)
  inputIndex=$(expr $ii % 3)
  modelIndex=$(expr $ii / 3)
  input=${inputArray[$inputIndex]}
  model=${modelArray[$modelIndex]}
  echo input:"$input" model:"$model"
done
```

Job array with a multiple nested indices, one numeric

You can frequently simplify things when one the lists you iterate over is numeric (e.g., each chromosome for each phenotype):

```
#!/bin/bash
...
#SBATCH --array=0-65
#SBATCH -J <jobname>
#SBATCH -o <output dir>/<jobname>_%a

ii=${SLURM_ARRAY_TASK_ID}
```

```
inputArray=(phenoA phenoB phenoC)
```

```
inputIndex=$(expr $ii / 22)
chrom=$(expr $(expr $ii % 22) + 1)
```

```
input=${inputArray[$inputIndex]}
```

```
ml load <modules>
```

```
program <args> \
  --phenotype "$input" \
  --genotypes <someprefix>_chr"$chrom"
```

Again, you can check that this works via a loop:

```
inputArray=(phenoA phenoB phenoC)

for ii in {0..65}
do
  inputIndex=$(expr $ii / 22)
  chrom=$(expr $(expr $ii % 22) + 1)
  input=${inputArray[$inputIndex]}
```

```
echo chrom:"$chrom" pheno:"$input"  
done
```

Useful links

- Batch scheduling and Summit specifications
- Blanca resources

Modifying this README

Do not directly modify this document directly. Instead modify `ibg_rc_cheatsheet.md` and run `compile.sh` to add TOC and create printable version. This will also update the README displayed on github. You might need to make `gh-md-toc` executable via `chmod a+x gh-md-toc` and you definitely will need to install pandoc. TOC generation powered by github-markdown-toc